

Virtual Machine Scale Sets (VMSS)

What is Scale Sets?

VM Scale Sets are intended to manage a group of identical, load-balanced VMs.

Automatically scales the number of VMs based on demand or schedule.

Ensures high availability and redundancy.



Key Features:

Autoscaling: Automatically scales VM instances on the basis of CPU, memory, or custom metrics.

Load Balancing: It integrates with Azure Load Balancer or Application Gateway.

Fault Tolerance: Spreads VMs across fault and update domains

Integration: Seamless integration with Azure Monitor and Azure DevOps.



Practical Example: **Deploying a Web App on Scale Sets.**

Create a Resource Group:

```
az group create  
--name ScaleSetRG  
--location "East US"
```

Create a Virtual Machine Scale Set:

```
az vmss create \  
--resource-group ScaleSetRG \  
--name MyScaleSet \  
--image UbuntuLTS \  
--admin-username azureuser \  
--generate-ssh-keys \  
--instance-count 2 \  
--upgrade-policy-mode Automatic \  
--load-balancer ""
```



Enable Autoscaling:

```
az monitor autoscale create \  
  --resource-group ScaleSetRG \  
  --resource MyScaleSet \  
  --resource-type  
Microsoft.Compute/virtualMachineScaleSets \  
  --name MyAutoScaleRule \  
  --min-count 2 \  
  --max-count 10 \  
  --count 3
```

Add a Scaling Rule: Scale out when CPU usage exceeds 70%:

```
az monitor autoscale rule create \  
  --resource-group ScaleSetRG \  
  --. autoscale-name MyAutoScaleRule \  
  --condition "Percentage CPU > 70 avg 5m" \  
  --scale out 1
```



Test the Setup:

Apply high traffic load on VMs by running a CPU intensive task.

Monitor scaling through the Azure Portal or CLI:

```
az vmss list-instances  
--resource-group ScaleSetRG  
--name MyScaleSet  
--output table
```

Access the Application:

Get the public IP of the Load Balancer:

```
az network public-ip show  
-resource-group ScaleSetRG  
--name MyScaleSetLBPublicIP  
--query ipAddress -o tsv
```

Open the IP address in a browser to view the hosted application.



2. Azure App Services

What are App Services?

A PaaS service offering of hosting web apps, REST APIs, and mobile backends.

A fully managed service with built-in scaling, security, and CI/CD capabilities.



Key Features

Language Support:

.NET, Java, Python, PHP, Node.js, and Ruby.

Autoscaling: Autoscale based on usage.

Custom Domains & SSL: Add secure connection with ease.

DevOps Integration: Integration with Azure DevOps, GitHub, and Bitbucket for CI/CD.

Deployment Slots: Test new versions in staging environments before production.



Practical Example: Deploying a Web App

Create a Resource Group:

```
az group create  
--name AppServiceRG  
--location "East US"
```

Create an App Service Plan:

The plan determines the compute resources for the app.

```
az appservice plan create \  
  --name MyAppServicePlan \  
  --resource-group AppServiceRG \  
  --sku B1 \  
  --is-linux
```

Create a Web App:

```
az webapp create \  
  --resource-group AppServiceRG \  
  --plan MyAppServicePlan \  
--name MyWebApp \  
  --runtime "PYTHON|3.8"
```



Deploy Your Code:

Use GitHub for CI/CD integration or deploy manually:

```
az webapp deployment source config \  
--name MyWebApp \  
--resource-group AppServiceRG \  
--repo-url https://github.com/<your-  
repo-url> \  
--branch main \  
--manual-integration
```

Access the Web App:

The default URL will be:
<https://<MyWebApp>.azurewebsites.net>



Enable Scaling:

```
az monitor autoscale create  
--resource-group AppServiceRG \  
--resource MyAppServicePlan \  
--resource-type Microsoft.Web/serverFarms \  
--name AppAutoScaleRule \  
--min-count 1 \  
--max-count 5 \  
--count 2
```

VM Scale Sets:

Use when you need full control over the underlying infrastructure, or for compute-heavy applications like data processing pipelines.

App Services:

Use for web apps, APIs, or lightweight applications where simplicity, scaling, and integration are the priority.

